

# Lab Week 7



Session: 2022 – 2026

## Submitted by:

Muhammad Saad Akmal

2022-CS-148

## Submitted To:

Sir Nauman Shafi

Department of Computer Science

**University of Engineering and Technology**

**Lahore Pakistan**

```

import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.datasets import fetch_california_housing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor

#Q1
iris = datasets.load_iris()
housing = fetch_california_housing()

iris_data = pd.DataFrame(iris.data , columns=iris.feature_names)
iris_target_columns = iris.target_names
iris_target = iris.target
iris_data_null = iris_data.isnull()

plt.scatter(iris_data['sepal length (cm)'], iris_data['sepal width (cm)'], c=iris_target)
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Scatter plot of Sepal Length vs Sepal Width')
plt.show()

housing_data = pd.DataFrame(housing.data , columns=housing.feature_names)
housing_target_columns = housing.target_names
housing_target = housing.target
housing_data_null = housing_data.isnull()

#Q2
iris_scaler = StandardScaler()
iris_scaler.fit(iris_data)
print(iris_scaler.transform(iris_data))

housing_scaler = StandardScaler()
housing_scaler.fit(housing_data)
print(housing_scaler.transform(housing_data))

iris_train_X, iris_test_X, iris_train_Y, iris_test_Y = train_test_split(iris_data , iris_target , random_state=104, test_size=0.20, shuffle=True)

fig, (train, test) = plt.subplots(1, 2)

train.scatter(iris_train_X['sepal length (cm)'], iris_train_X['sepal width (cm)'], c=iris_train_Y)
train.set_xlabel('Sepal Length (cm)')
train.set_ylabel('Sepal Width (cm)')
train.set_title('Scatter plot of Sepal Length vs Sepal Width (Training set)')

test.scatter(iris_test_X['sepal length (cm)'], iris_test_X['sepal width (cm)'], c=iris_test_Y)
test.set_xlabel('Sepal Length (cm)')
test.set_ylabel('Sepal Width (cm)')
test.set_title('Scatter plot of Sepal Length vs Sepal Width (Testing set)')

plt.show()

housing_train_X, housing_test_X, housing_train_Y, housing_test_Y = train_test_split(housing_data , housing_target , random_state=104, test_size=0.20, shuffle=True)

#TASK2

#Q3
accuracy_list = []
fig, axes = plt.subplots(2, 2, figsize=(8, 15))
for i in range(1,16):
    neigh = KNeighborsClassifier(n_neighbors=i)
    neigh.fit(iris_train_X , iris_train_Y)
    predicted = neigh.predict(iris_test_X)
    accuracy = metrics.accuracy_score(iris_test_Y, predicted)
    accuracy_list.append(accuracy)
    print(f'For iteration {i} the accuracy is {accuracy}')

axes[0,0].plot([i for i in range(1,16)] , accuracy_list)
axes[0,0].set_xlabel('K')
axes[0,0].set_ylabel('Accuracy')
axes[0,0].set_title('Accuracy of KNN')

```

```

#Q4
linear = SVC(kernel='linear').fit(iris_train_X, iris_train_Y)
linear_predict = linear.predict(iris_test_X)
linear_accuracy = metrics.accuracy_score(iris_test_Y, linear_predict)

poly = SVC(kernel='poly').fit(iris_train_X, iris_train_Y)
poly_predict = poly.predict(iris_test_X)
poly_accuracy = metrics.accuracy_score(iris_test_Y, poly_predict)

rbf = SVC(kernel='rbf').fit(iris_train_X, iris_train_Y)
rbf_predict = rbf.predict(iris_test_X)
rbf_accuracy = metrics.accuracy_score(iris_test_Y, rbf_predict)

kernels = ['Linear', 'Poly', 'RBF']
svc_accuracy = [linear_accuracy, poly_accuracy, rbf_accuracy]
axes[0,1].bar(kernels, svc_accuracy, color=['Red', 'Blue', 'Green'])
axes[0,1].set_xlabel('Kernels')
axes[0,1].set_ylabel('Accuracy')
axes[0,1].set_title('Accuracy of SVC with different kernels')

#Q5
calssifier_10 = RandomForestClassifier(n_estimators = 10).fit(iris_train_X, iris_train_Y)
predict_10 = calssifier_10.predict(iris_test_X)
acc_10 = metrics.accuracy_score(iris_test_Y, predict_10)
print(f'Accuracy for estimator 10 {acc_10}')

calssifier_50 = RandomForestClassifier(n_estimators = 50).fit(iris_train_X, iris_train_Y)
predict_50 = calssifier_50.predict(iris_test_X)
acc_50 = metrics.accuracy_score(iris_test_Y, predict_50)
print(f'Accuracy for estimator 50 {acc_50}')

#TASK 3

regr = LinearRegression()

regr.fit(housing_train_X, housing_train_Y)
pred = regr.predict(housing_test_X)

mean_square_error = mean_squared_error(housing_test_Y, pred)
r2_Score = r2_score(housing_test_Y, pred)

print(f'Mean square error = {mean_square_error}')
print(f'R2 score = {r2_Score}')

regressor = DecisionTreeRegressor(random_state = 0)

regressor.fit(housing_train_X, housing_train_Y)
tree_pred = regressor.predict(housing_test_X)

tree_mean_square_error = mean_squared_error(housing_test_Y, tree_pred)
tree_r2_Score = r2_score(housing_test_Y, tree_pred)
print(f'Mean square error = {tree_mean_square_error}')
print(f'R2 score = {tree_r2_Score}')

indices = range(len(housing_test_Y))

fig, axes = plt.subplots(1, 2)

axes[0].scatter(indices, housing_test_Y, color='blue', alpha=0.6, label='Actual Values')
axes[0].scatter(indices, pred, color='red', alpha=0.6, label='Predicted Values')
axes[0].set_xlabel('Sample Index')
axes[0].set_ylabel('Value')
axes[0].set_title('Actual and Predicted Values (Linear Regression)')
axes[0].legend()
axes[0].grid()

regressor = DecisionTreeRegressor(random_state=0)
regressor.fit(housing_train_X, housing_train_Y)
tree_pred = regressor.predict(housing_test_X)

axes[1].scatter(indices, housing_test_Y, color='blue', alpha=0.6, label='Actual Values')
axes[1].scatter(indices, tree_pred, color='red', alpha=0.6, label='Predicted Values')
axes[1].set_xlabel('Sample Index')
axes[1].set_ylabel('Value')
axes[1].set_title('Actual and Predicted Values (Decision Tree)')
axes[1].legend()
axes[1].grid()

plt.tight_layout()
plt.show()

```

```

regressor = DecisionTreeRegressor(random_state=0)
regressor.fit(housing_train_X, housing_train_Y)
tree_pred = regressor.predict(housing_test_X)

axes[1].scatter(indices, housing_test_Y, color='blue', alpha=0.6, label='Actual Values')
axes[1].scatter(indices, tree_pred, color='red', alpha=0.6, label='Predicted Values')
axes[1].set_xlabel('Sample Index')
axes[1].set_ylabel('Value')
axes[1].set_title('Actual and Predicted Values (Decision Tree)')
axes[1].legend()
axes[1].grid()

plt.tight_layout()
plt.show()

model = ['LinearRegression', 'Decision Tree']
mse = [mean_square_error, tree_mean_square_error]
r2 = [r2_Score, tree_r2_Score]

fig, axes = plt.subplots(1, 2)

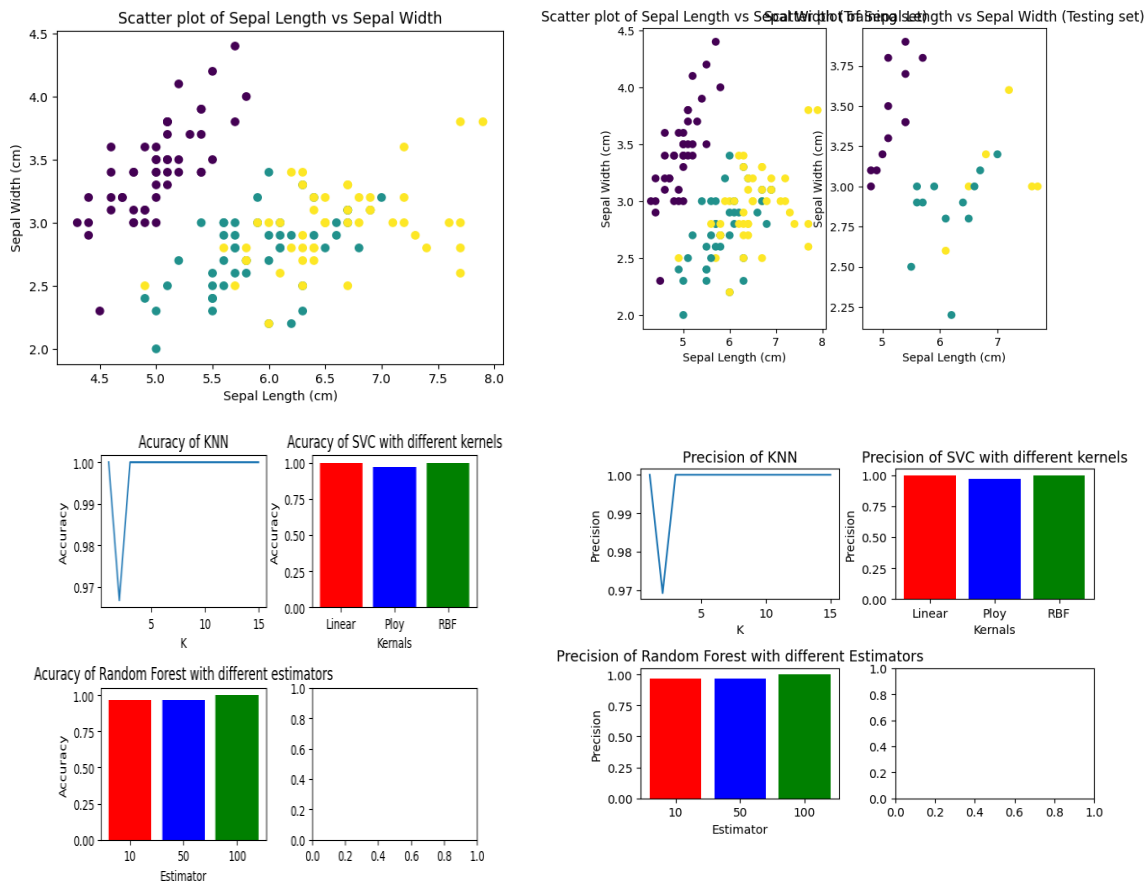
axes[0].bar(model, mse, color=['Red', 'Blue'])
axes[0].set_xlabel('Model')
axes[0].set_ylabel('Mean Squared Error (MSE)')
axes[0].set_title('MSE of Models')









# Plot R2 Score
axes[1].bar(model, r2, color=['Red', 'Blue'])
axes[1].set_xlabel('Model')
axes[1].set_ylabel('R2 Score')
axes[1].set_title('R2 Score of Models')

plt.show()


```

Output:





Actual and Predicted Values (Linear Regression)



Actual and Predicted Values (Decision Tree)

